

# Race conditions

ID2010

2022

# Race condition

- IF multiple processes or threads compete for a non-sharable and mutable resource
- AND access to the resource is not controlled
- THEN there is a *race* to modify the resource

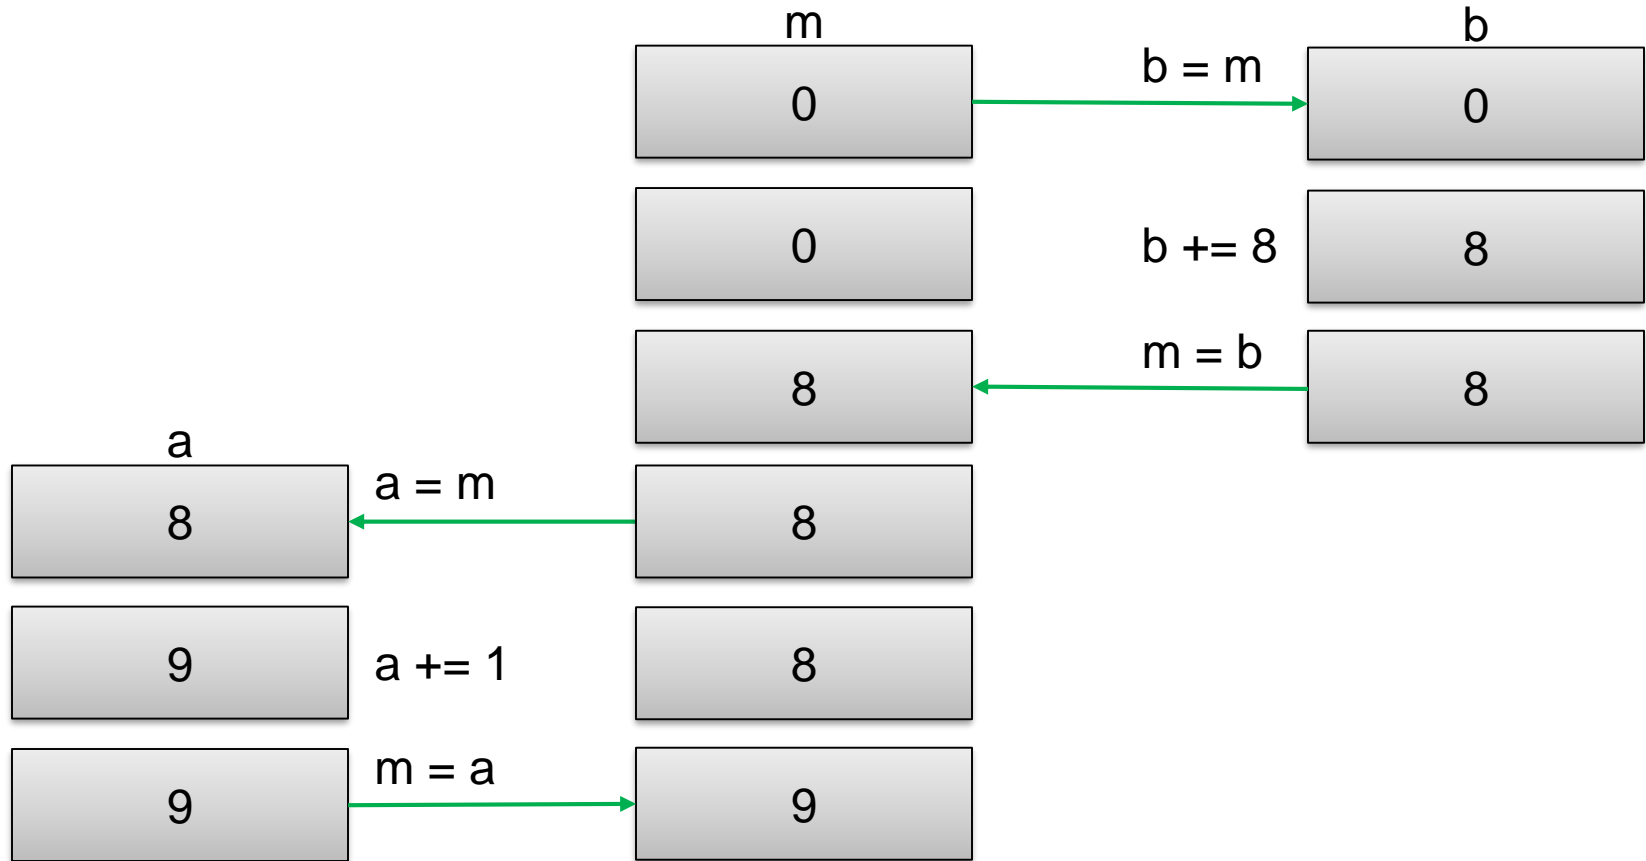
# Race condition

- Mutable (writable) resources:
  - register, memory, file, hardware port
- Non-shared access:
  - Read-Modify-Write
- Shared and uncontrolled access:
  - Read-Modify-**Read**-Write-**Modify**-**Write**

# Race condition


- State is corrupted by overlapping updates
- The race condition means that a corruption is *possible*
- The probability of a corruption actually happening depends on the system
- non-deterministic errors, bugs are hard to find

# Race condition on a variable




Thread a and thread b updates m separately.  
The value of m is 9 (correct).

# Race condition on a variable



The update from thread b is overwritten by thread a.  
The value of m is 1, not 9 (wrong by 8).


# Race condition on a variable



The update from thread a is overwritten by thread b.  
The value of m is 8, not 9 (wrong by 1).

# Race condition on a variable


- The race condition is a possible bug
- The bug only happens when a and b overlap
- Errors in variable m accumulate over time  
and may not be recognized at first:  
 $0 + 1 + 0 + 8 + 1 + 1 + 0 + 8 + \dots$





# Generalizations


- A single variable is a special case
- of a block of data that must be internally and externally *consistent*




# History

- 'Race condition' was used in 1954 referring to signals in electrical circuits
- Logic gates in particular may suffer if signals arrive at different times
- In software also known as 'data race'

# Non-serialized operations



# Serialized operations



The resource is protected by a mechanism which enforces serial access to it:

- Operating system
- Database manager
- Programming language synchronization

# Synchronization

- Only one process may execute the critical section of code
- Acquire exclusive rights
- Execute critical section
- Release exclusive rights

# Synchronization

```
...  
acquire_permission(s) ;  
a = m ;  
a = a + k ;  
m = a ;  
release_permission(s) ;  
...
```


The entity **s** controls access to **m**.

# Synchronization

- Semaphore – guards  $n$  resources
  - Any thread can return a resource
- Mutex – guards 1 resource
  - Only the currently owning thread can return the resource
- Threads block until the resource is granted

# Semaphore

- *Counts the no of available resources*





# Semaphore

- *Counts the no of available resources*



1. Decrement counter

# Semaphore

- *Counts the no of available resources*



1. Decrement counter
2. Use resource

# Semaphore


- *Counts the no of available resources*



1. Decrement counter
2. Use resource
3. Increment counter

# Semaphore

- *Counts the no of available resources*




*No identification of resources.*

1. Decrement counter
2. Use resource
3. Increment resource

# Semaphore

- *Counts the no of available resources*



1. Decrement counter
2. Use resource
3. Increment resource


*No identification of resources.  
Increment and decrement are atomic.*

# Semaphore caveats

- Processes (i.e. programmers) must follow the protocol and **not**:
  - Forget to return a resource after use
  - Return a resource that was not requested
  - Hold a resource for too long
  - Use a resource anyway

# Mutex

- A binary semaphore (one resource)



# Mutex

- A binary semaphore (one resource)


1. Acquire





# Mutex

- A binary semaphore (one resource)



1. Acquire
2. Use

# Mutex

- A binary semaphore (one resource)

1. Acquire
2. Use
3. Return



# Mutex


- A binary semaphore (one resource)
- The process acquiring the resource is the only one that can return it (ownership)
- The *synchronized* keyword in Java use any object as a monitor (a kind of mutex).

# Monitor

- A mutex with a thread queue
- Threads queue for the mutex
- Threads can yield the resource (wait)
- Threads can alert other threads (notify)
- In Java, every object has a monitor
- The *synchronized* keyword

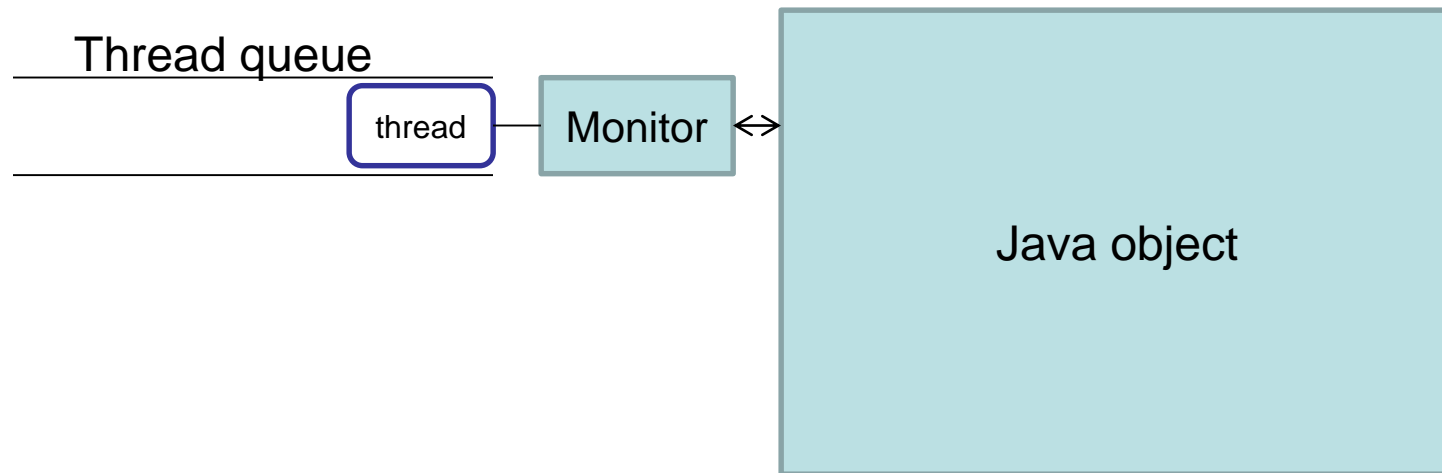
# Monitor

- A mutex with a thread queue



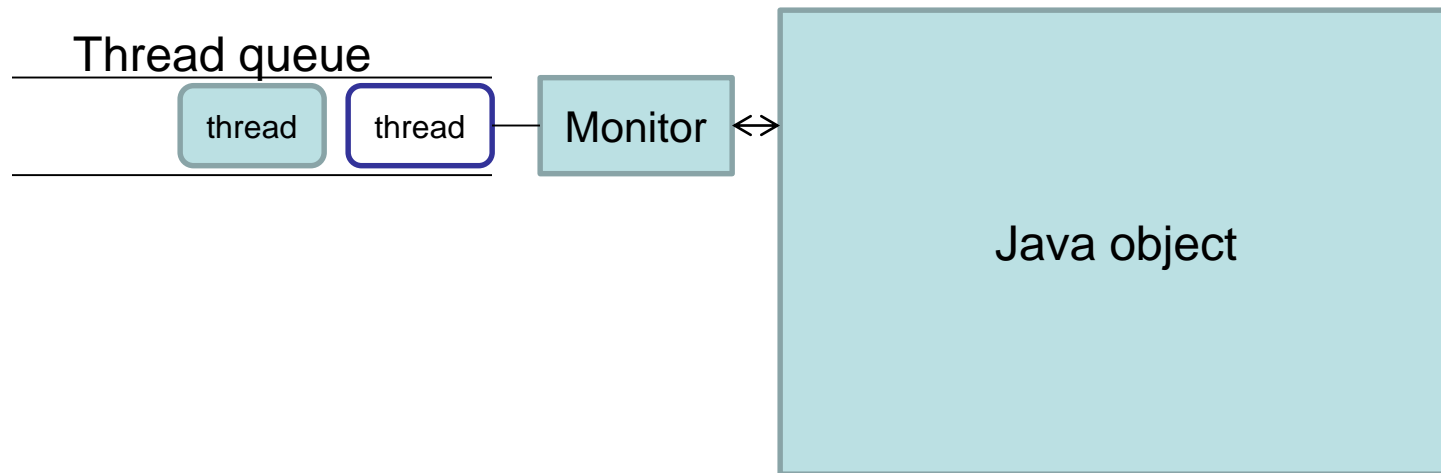
# Monitor

- A mutex with a thread queue



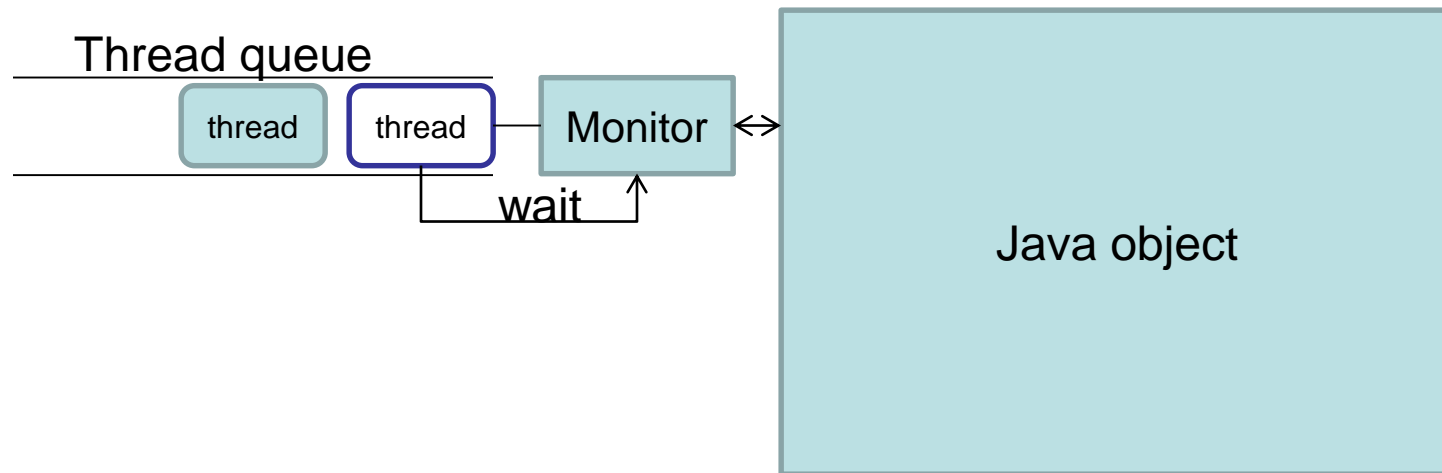
# Monitor

- A mutex with a thread queue



# Monitor


- A mutex with a thread queue





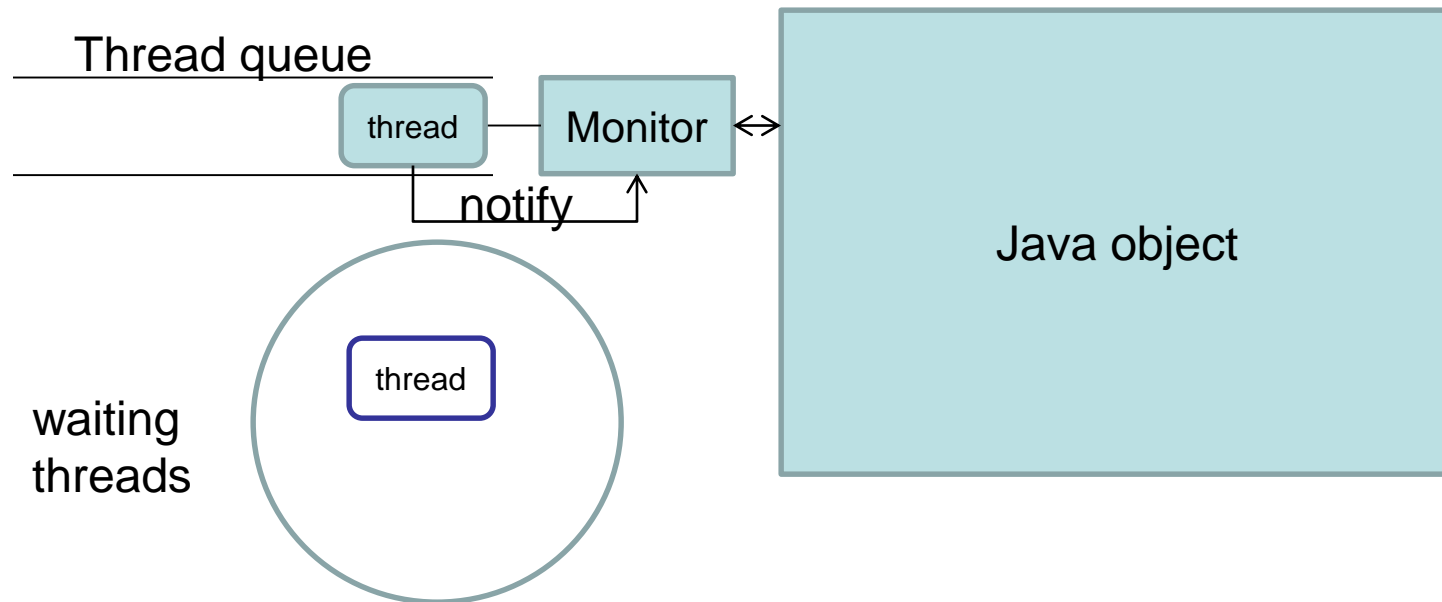
# Monitor

- A mutex with a thread queue




# Monitor

- A mutex with a thread queue




# Monitor

- A mutex with a thread queue



# Monitor

- A mutex with a thread queue



# Some unsharable resources

- Process memory (part of)
- Firmware memory (NVRAM, EEPROM)
- Filesystems
- Databases
- Communication lines
- Network adapters
- User interface devices

# Lock-free code


- Semaphores, mutexes, and monitors may block (suspend) the thread
- The thread is locked in waiting
- The thread now depends on other threads:
  - not being locked in turn
  - not being delayed
  - being bug-free

# Lock-free code

- Lock-free code, non-blocking code

```
while true {  
  
    previousState = getCurrentState()  
  
    nextState = previousState + operation  
  
    if atomicCompareAndSet(previousState, nextState)  
        break  
  
    // else try again  
}
```

# Lock-free code





# Lock-free code


- Complicated and inefficient algorithms
- The ABA problem:
  - thread X reads state A
  - while X is suspended, thread Y removes A, puts on B, removes B, puts back a different A
  - thread X updates since  $A==A$ , but they are in fact not the same A

# Example ABALinkedList

Thread 1

Pop the top number off  
and save it

Stack



Thread 2

Pop top – 1 off and save it:  
a = pop()  
b = pop()  
push(a)  
save b

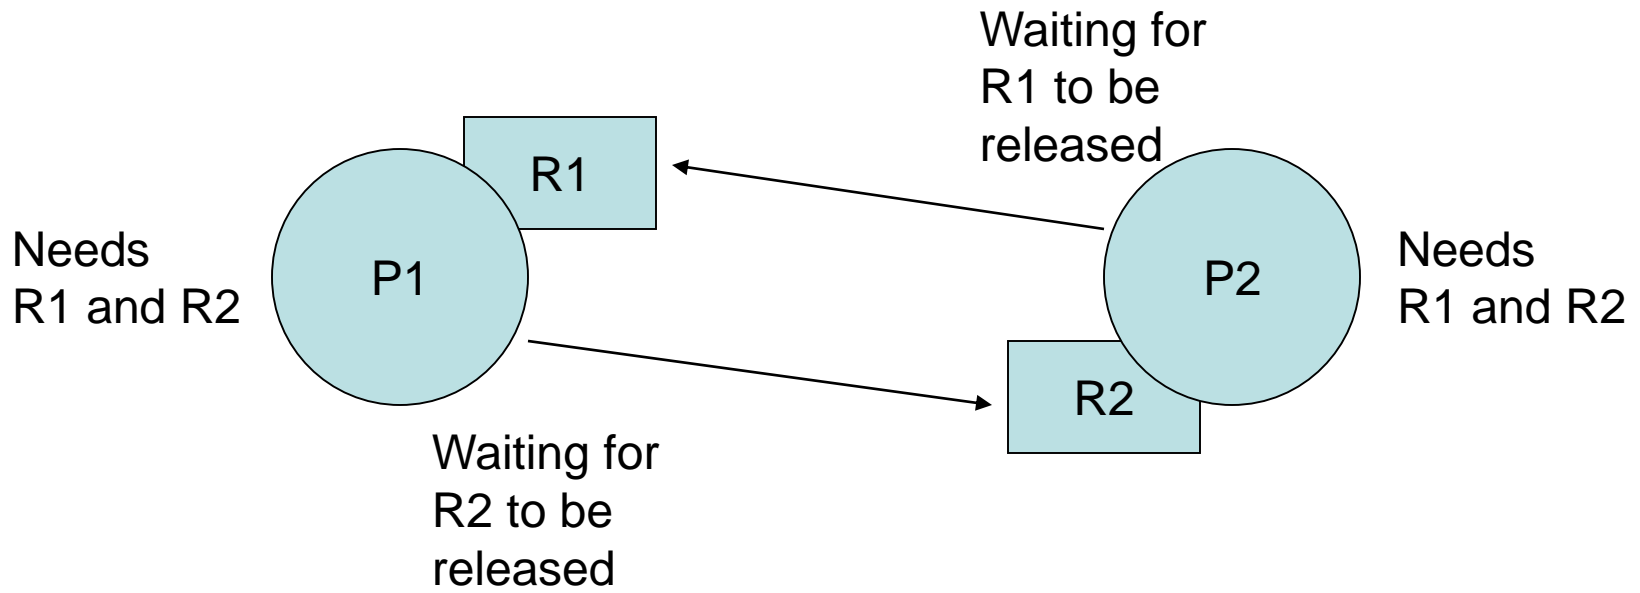
## Consistency assumption:

When the stack is empty,  
every number originally on the stack  
is found exactly once in either  
thread 1 or in thread 2.

0

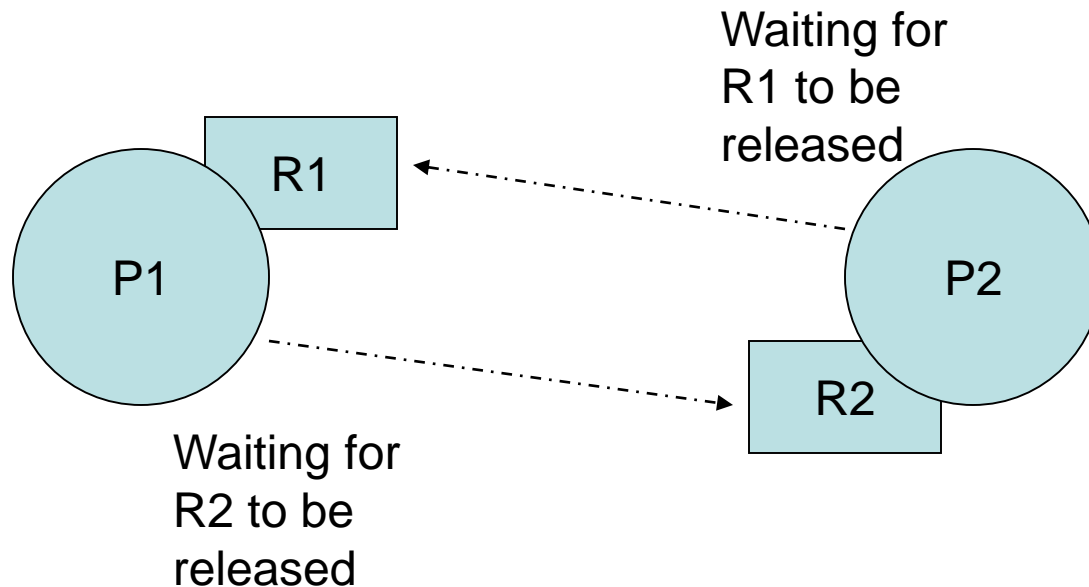
# Deadlock

- Processes wait forever for all required resources to be released



# Livelock

- Processes attempt to break the deadlock by time-out and release, but no-one wins



# Famous race conditions

- The Therac-25 radiation therapy machine (1987)
  - Safety flag was not set but incremented (wrap) + fast manual input outran proper radiation shielding
- Northeast US electrical blackout (2003)
  - Alert message systems broke down; human operators unaware of critical situation – simultaneous write access: alert sys loop.
- Citibank withdrawal fraud (2012)
  - Multiple withdrawals from ATMs within a 60-second window did not accumulate debit on the account

# Further fun reading

- Dijkstra, The Dining Philosophers problem
- Chandy/Misra solution to DPh problem
- Producer-consumer problem
- Sleeping barber problem
- Readers-writers problem
- Cigarette smokers problem
- ABA problem

End